

# **ECE 448/528**

## **Application Software Design**

### **Lecture 7. TCP Server Design II**

#### **Spring 2025**

**Won-Jae Yi, Ph.D.**

**Department of Electrical and Computer Engineering**  
**Illinois Institute of Technology**

# Implementing ReverseProcessor

# Handling Multiple Strings

- Recall telnet separates lines by `"\r\n"`.
- User story to support multiple strings:
  - As an end-user, I want to send multiple strings separated by `"\r\n"` to the server to be reversed via UTF-8 encoding, so that I can use telnet to test the server.
- User story to support end of request:
  - As an end-user, I want to send an empty string to the server to indicate I am done, so that the connection can be closed gracefully.
- Which one to implement first?
  - Take the simpler one!

# Obtain String and Disconnect if Requested

- `src/main/ece448/lec06/ReverseProcessor.java`
  - Under branch `lec07-message`.
- Use `ByteArrayOutputStream` to buffer bytes before they are converted to a `String`
- Ignore `\r` since that's not part of the string.
- When `\n` arrives, extract the string via `toString`.
  - Decode from UTF-8.
  - Exit if the string is empty.
  - Log and continue to the next string.
  - Don't forget to reset the buffer!

# Reverse and Send Back

- `src/main/ece448/lec06/ReverseProcessor.java`
  - Under branch `lec07-reverse`.
- Call `ReverseString.reverse` to reverse the string.
- Call `getBytes("UTF-8")` to get the bytes.
- Send the bytes back using `OutputStream.write`.
- Send `"\r\n"` so that telnet will display the reversed string in its line.

# Discussions

# General TCP Server Architecture

- Create a server socket to listen on a port, then use a loop to wait/accept new clients.
- Once a client arrives, send the client socket to a separated thread to serve the client. Then, the server socket may continue accepting additional clients.
- To process requests from a client, use another loop to read from the input byte stream.
  - When a complete message is received, the client will perform some computation to obtain a result.
  - The result will be sent back to the client via the output byte stream.
- Need a presentation layer protocol to decide
  - How to extract messages from the input byte stream.
  - How to convert results to the output byte stream.

# Design Considerations

- The two responsibilities of waiting for clients and serving a client can be handled by two classes.
- If a different kind of TCP service should be implemented, the class that waits for clients to connect will mostly remain the same.
- However, we may need to rewrite the class that serves a client.
- “Sandwich” structure to process requests from clients
  - Decode a request message from the input TCP stream.
  - Do some computations to obtain a response.
  - Encode and send the response via the output TCP stream.
- A well-designed framework may allow us to focus on the application layer only.
  - To only consider how to compute responses from requests, without worrying about how they are decoded/encoded or sent/received over the network.



# Performance Considerations

- It is not efficient to read one byte at a time from the input TCP stream.
  - Prefer to read all available bytes at a time – but this requires additional code to handle the buffer properly.
- It is “expensive” to create and maintain a thread for each client.
  - A major problem when you need to serve thousands or more clients at the same time.
  - OS needs to allocate memory for each thread.
  - Frequent context switches consume a lot of processor power.
  - Prefer to use thread pools or I/O multiplexing techniques that use less threads – but this requires substantial change to the whole design.
- We would move to a well-designed framework for our projects to meet the performance demands you may face in the future.

# General TCP Server Architecture

- ▶ Create a server socket to listen on a port, and then use a loop to wait/accept new clients.
- ▶ Once a client arrives, send the client socket to a separated thread to serve the client. Then the server socket may continue accepting additional clients.
- ▶ To process requests from a client, use another loop to read from the input byte stream.
  - ▶ When a complete message is received, the client will perform some computation to obtain a result.
  - ▶ The result will be sent back to the client via the output byte stream.
- ▶ Need a presentation layer protocol to decide
  - ▶ How to extract messages from the input byte stream.
  - ▶ How to convert results to the output byte stream.

# Design Considerations

- ▶ The two responsibilities of waiting for clients and serving a client can be handled by two classes.
- ▶ If a different kind of TCP service should be implemented, the class that waits for clients to connect will mostly remain the same.
- ▶ However, we may need to rewrite the class that serves a client.
- ▶ “Sandwich” structure to process requests from clients
  - ▶ Decode a request message from the input TCP stream.
  - ▶ Do some computation to obtain a response.
  - ▶ Encode and send the response via the output TCP stream.
- ▶ A well-designed framework may allow us to focus on the application layer only.
  - ▶ To only consider how to compute responses from requests, without worrying about how they are decoded/encoded or sent/received over the network.

# Performance Considerations

- ▶ It is not efficient to read one byte at a time from the input TCP stream.
  - ▶ Prefer to read all available bytes at a time – but this requires additional code to handle the buffer properly.
- ▶ It is costly to create and maintain a thread for each client.
  - ▶ A major problem when you need to serve thousands or more clients at the same time.
  - ▶ OS need to allocate memory for each thread.
  - ▶ Frequent context switches consume a lot of processor power.
  - ▶ Prefer to use thread pools or I/O multiplexing techniques that use less threads – but this requires substantial change to the whole design.
- ▶ We would move to a well-designed framework for our projects to meet the performance demands you may face in future.